


Adversarial examples

Overview

- What are adversarial examples?
- Why do they happen?
- How can they be used to compromise machine learning systems?
- What are the defenses?
- How to use adversarial examples to improve machine learning, even when there is no adversary

Since 2013, deep neural networks have matched human performance at...



(Szegedy et al, 2014)

...recognizing objects and faces...



(Taigmen et al, 2013)



(Goodfellow et al, 2013)


...solving CAPTCHAS and reading addresses...



(Goodfellow et al, 2013)

and other tasks...

Adversarial Examples



Timeline:

“Adversarial Classification” Dalvi et al 2004: fool spam filter

“Evasion Attacks Against Machine Learning at Test Time”

Biggio 2013: fool neural nets


Szegedy et al 2013: fool ImageNet classifiers imperceptibly

Goodfellow et al 2014: cheap, closed form attack

Turning Objects into “Airplanes”




Attacking a Linear Model




Not just for neural nets

- Linear models
 - Logistic regression
 - Softmax regression
 - SVMs
- Decision trees
- Nearest neighbors

Adversarial Examples from Overfitting




Adversarial Examples from Excessive Linearity




Modern deep nets are very piecewise linear


Rectified linear unit




Maxout




Carefully tuned sigmoid



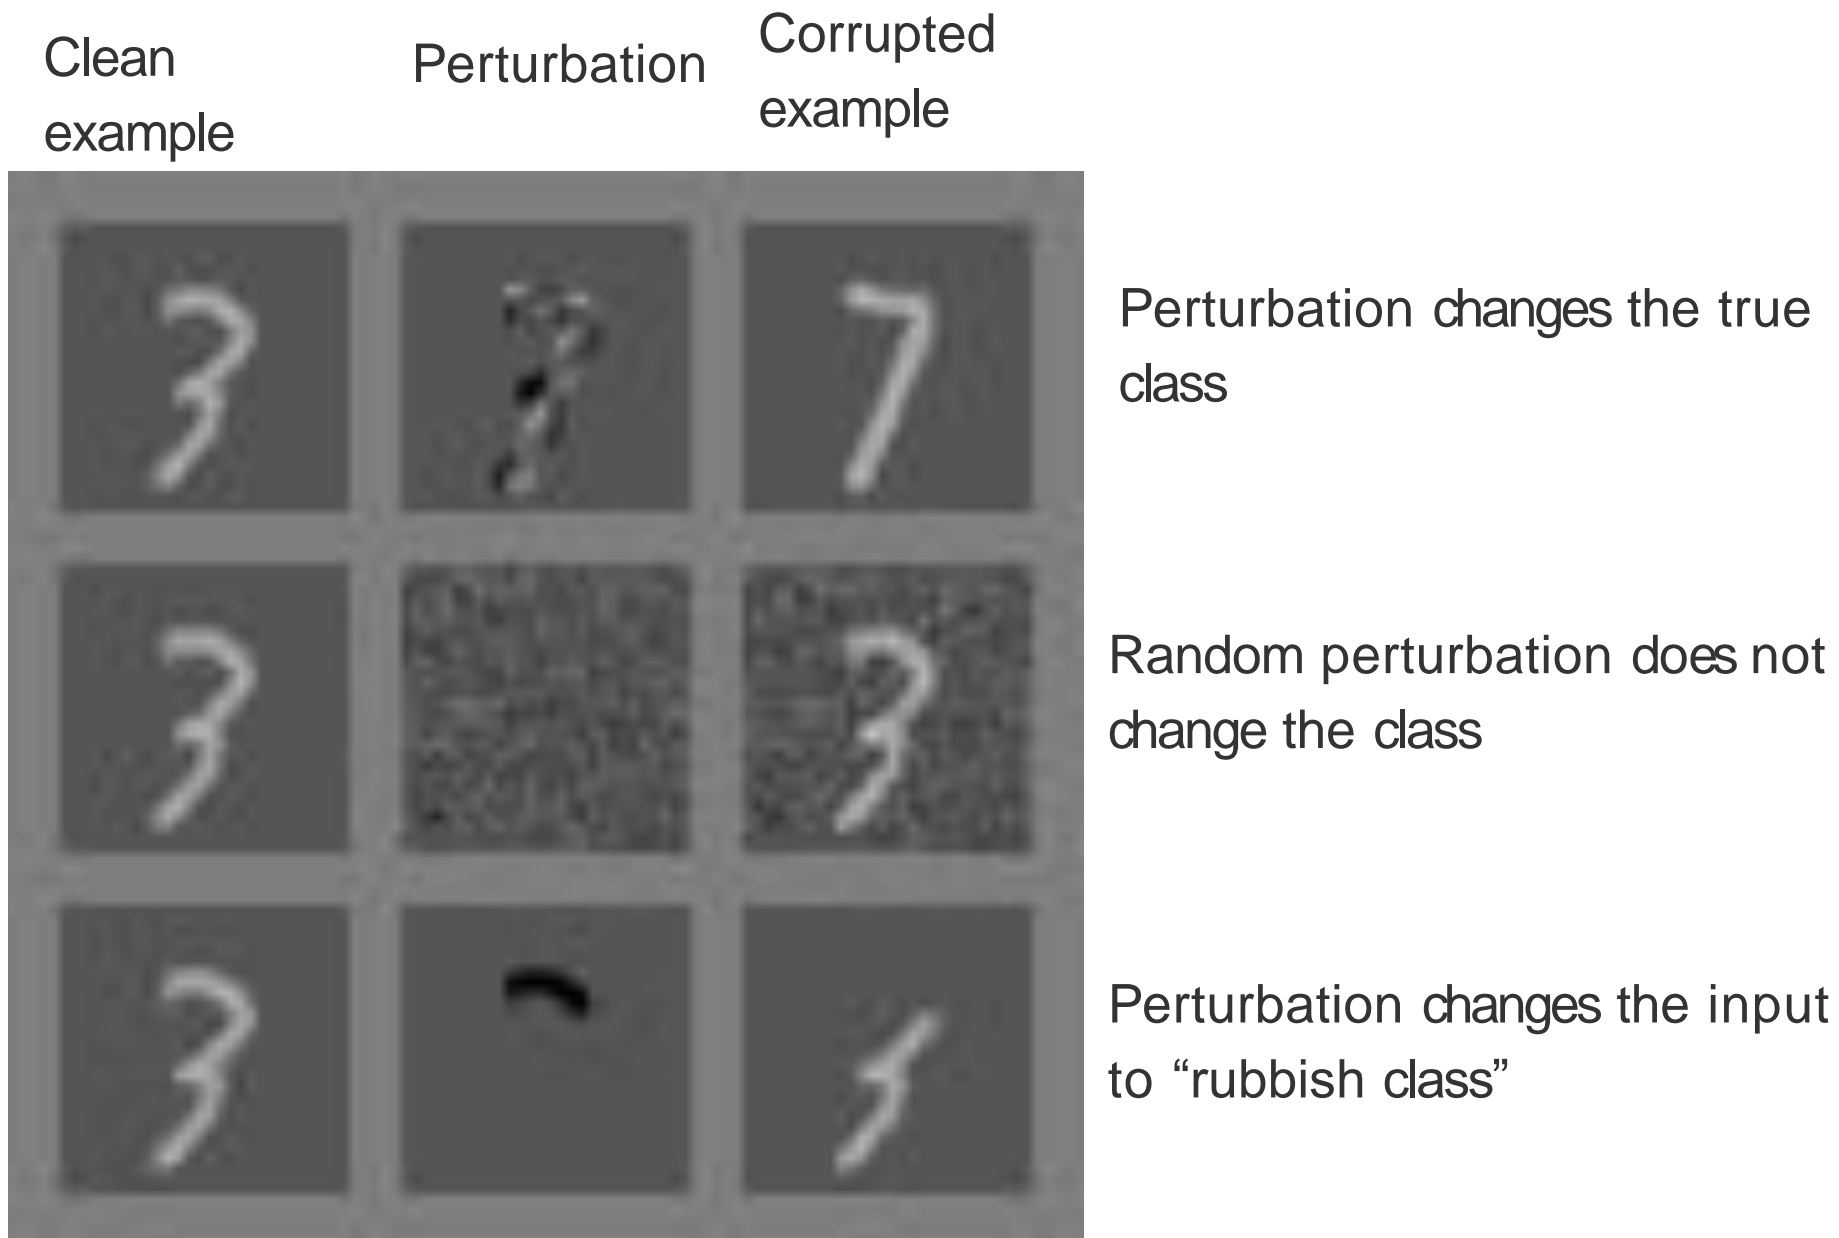
LSTM



Nearly Linear Responses in Practice



Small inter-class distances



All three perturbations have L2 norm 3.96

This is actually small. We typically use 7!

The Fast Gradient Sign Method

$$J(\tilde{\mathbf{x}}, \boldsymbol{\theta}) \approx J(\mathbf{x}, \boldsymbol{\theta}) + (\tilde{\mathbf{x}} - \mathbf{x})^\top \nabla_{\mathbf{x}} J(\mathbf{x}).$$

Maximize



$$J(\mathbf{x}, \boldsymbol{\theta}) + (\tilde{\mathbf{x}} - \mathbf{x})^\top \nabla_{\mathbf{x}} J(\mathbf{x})$$

subject to


$$\|\tilde{\mathbf{x}} - \mathbf{x}\|_\infty \leq \epsilon \quad (\text{maxnorm})$$

$$\Rightarrow \tilde{\mathbf{x}} = \mathbf{x} + \epsilon \text{sign}(\nabla_{\mathbf{x}} J(\mathbf{x})).$$

Maps of Adversarial and Random Cross-Sections





Maps of Adversarial Cross-Sections




Maps of Random Cross-Sections

Adversarial examples
are not noise




Estimating the Subspace Dimensionality




(Tramèr et al, 2017)

Wrong almost everywhere



High-Dimensional Linear Models


Weights




Signs of weights




Clean examples



Adversarial




Linear Models of ImageNet




(Andrej Karpathy, "Breaking Linear Classifiers on ImageNet")

RBFs behave more intuitively



Cross-model, cross-dataset generalization



Cross-technique transferability

| Source Machine Learning Technique | DNN | LR | SVM | DT | kNN | Ens. |
|-----------------------------------|-------|-------|-------|-------|-------|-------|
| DNN | 38.27 | 23.02 | 64.32 | 79.31 | 8.36 | 20.72 |
| LR | 6.31 | 91.64 | 91.43 | 87.42 | 11.29 | 44.14 |
| SVM | 2.51 | 36.56 | 100.0 | 80.03 | 5.19 | 15.67 |
| DT | 0.82 | 12.22 | 8.85 | 89.29 | 3.31 | 5.11 |
| kNN | 11.75 | 42.89 | 82.16 | 82.95 | 41.65 | 31.92 |
| Target Machine Learning Technique | DNN | LR | SVM | DT | kNN | Ens. |

(Papernot 2016)

Transferability Attack

Target model with unknown weights, machine learning algorithm, training set; maybe non-differentiable

Train your own model

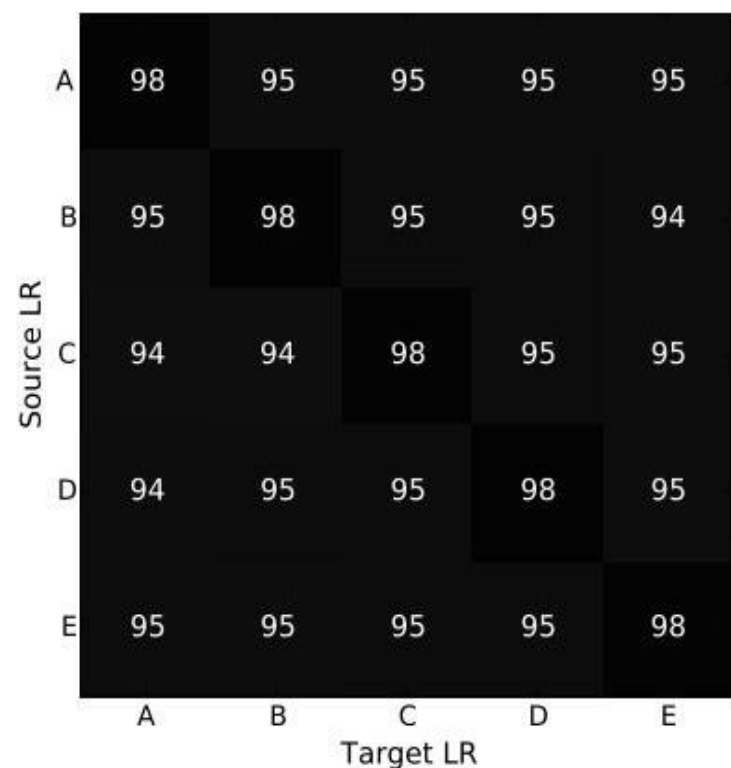
Substitute model mimicking target model with known, differentiable function

Adversarial crafting against substitute


Adversarial examples

Deploy adversarial examples against the target; transferability property results in them succeeding

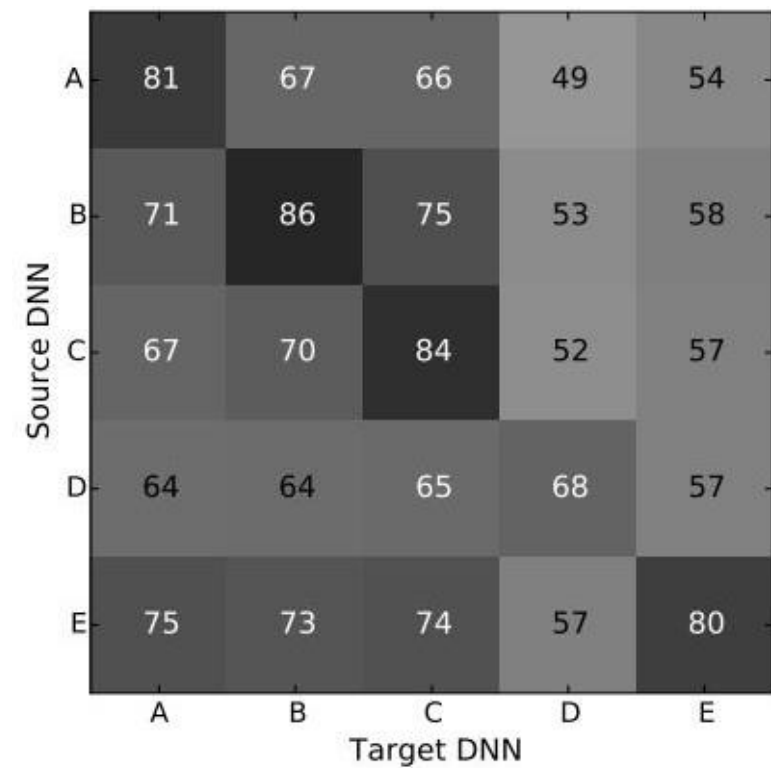
Cross-Training Data Transferability



Strong



Weak



Intermediate

(Papernot 2016)


Enhancing Transfer With Ensembles

| | RMSD | ResNet-152 | ResNet-101 | ResNet-50 | VGG-16 | GoogLeNet |
|-------------|-------|------------|------------|-----------|--------|-----------|
| -ResNet-152 | 17.17 | 0% | 0% | 0% | 0% | 0% |
| -ResNet-101 | 17.25 | 0% | 1% | 0% | 0% | 0% |
| -ResNet-50 | 17.25 | 0% | 0% | 2% | 0% | 0% |
| -VGG-16 | 17.80 | 0% | 0% | 0% | 6% | 0% |
| -GoogLeNet | 17.41 | 0% | 0% | 0% | 0% | 5% |

Table 4: Accuracy of non-targeted adversarial images generated using the optimization-based approach. The first column indicates the average RMSD of the generated adversarial images. Cell (i, j) corresponds to the accuracy of the attack generated using four models except model i (row) when evaluated over model j (column). In each row, the minus sign “-” indicates that the model of the row is not used when generating the attacks. Results of top-5 accuracy can be found in the appendix (Table 14).

(Liu et al, 2016)

Adversarial Examples in the Human Brain




These are
concentric
circles,
not
intertwined
spirals.

(Pinna and Gregory, 2002)

Practical Attacks

- Fool real classifiers trained by remotely hosted API (MetaMind, Amazon, Google)
- Fool malware detector networks
- Display adversarial examples in the physical world and fool machine learning systems that perceive them through a camera

Adversarial Examples in the Physical World



(Kurakin et al, 2016)

Failed defenses

Generative
pretraining

Removing perturbation
with an autoencoder

Adding noise
at test time

Ensembles

Confidence-reducing
perturbation at test time

Error correcting
codes

Multiple glimpses

Weight decay


Adding noise
at train time

Double backprop

Various
non-linear units

Dropout

Training on Adversarial Examples



Adversarial Training

Labeled as bird



Still has same label (bird)



Decrease
probability
of bird class

Virtual Adversarial Training

Unlabeled; model
guesses it's probably
a bird, maybe a plane




New guess should
match old guess
(probably bird, maybe plane)



Adversarial
perturbation
intended to
change the guess

Text Classification with VAT

RCV1 Misclassification Rate



Zoomed in for legibility

Universal engineering machine (model-based optimization)

Make new inventions
by finding input that
maximizes model's
predicted
performance

Training data

Extrapolation



Conclusion

- Attacking is easy
- Defending is difficult
- Adversarial training provides regularization and semi-supervised learning
- The out-of-domain input problem is a bottleneck for model-based optimization generally